Introduction to Data Science and IBM's Data Science Experience

Power of data. Simplicity of design. Speed of innovation.

Joel Patterson
Bernie Beekman
Davin Shearer

The digital age is changing the way we live, play, learn and work…

Google

Google Search I'm Feeling Lucky

What is the Data Scientist?

Scripting, SQL
Python, R Scala
Data Pipelines
Big Data/ Apache Spark

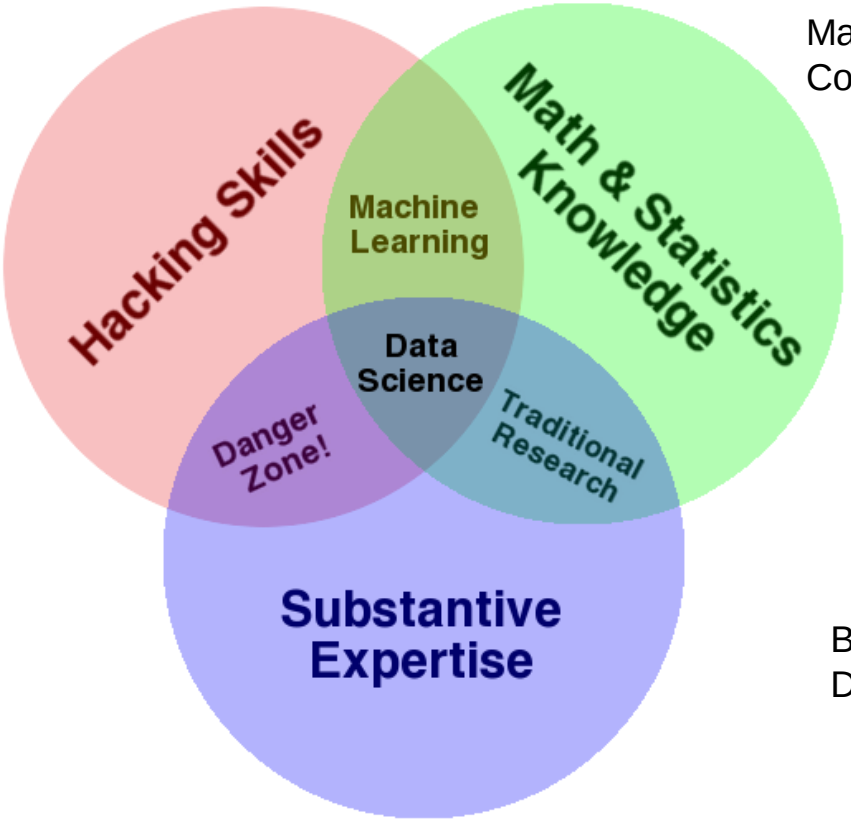

Mathematical Background
Computational Science

Business/Industry Expertise
Domain Knowledge
- Supply Chain
- CRM
- Financials
- Networking

Drew Conway's Data Science Venn Diagram

Google Trends – Data Science Languages

SPSS

SAS

Python

R

Scala

Trends in Google Searches (September 2nd 2016)

The perfect Data Science Team

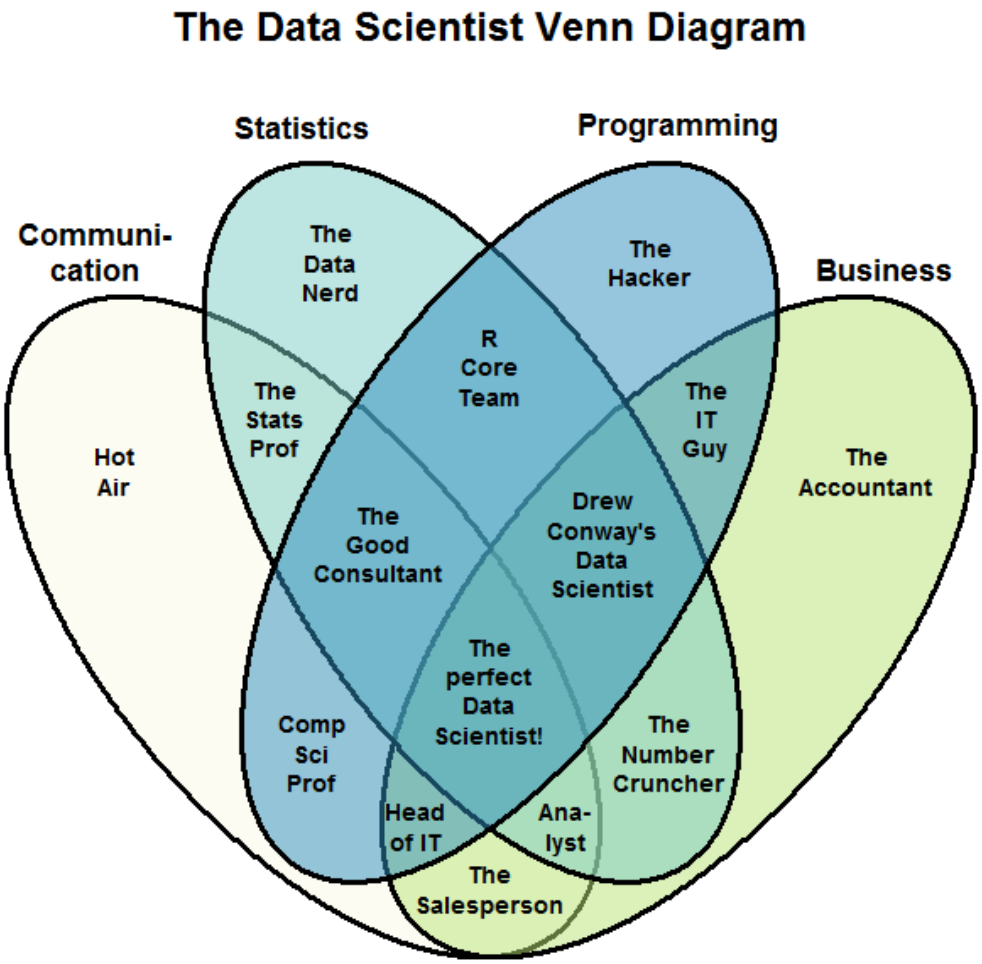

Normally not all the skills are in one single person but rather in a data science team
In IBM Data Science Experience we include tools to make the perfect Data Science Team
All in a collaborative, cloud environment that scales in demand

IBM Watson Data Platform

Mission: Make Data Simple and Accessible to All

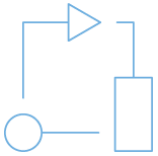

Platform.

Method.

Ecosystem.

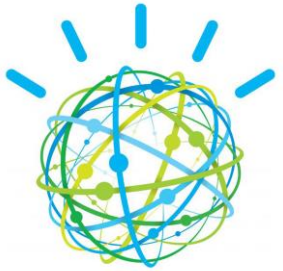

http://ibm.co/makedatasimple

IBM Watson Data Platform

Experience New Ways To Put Data To Work

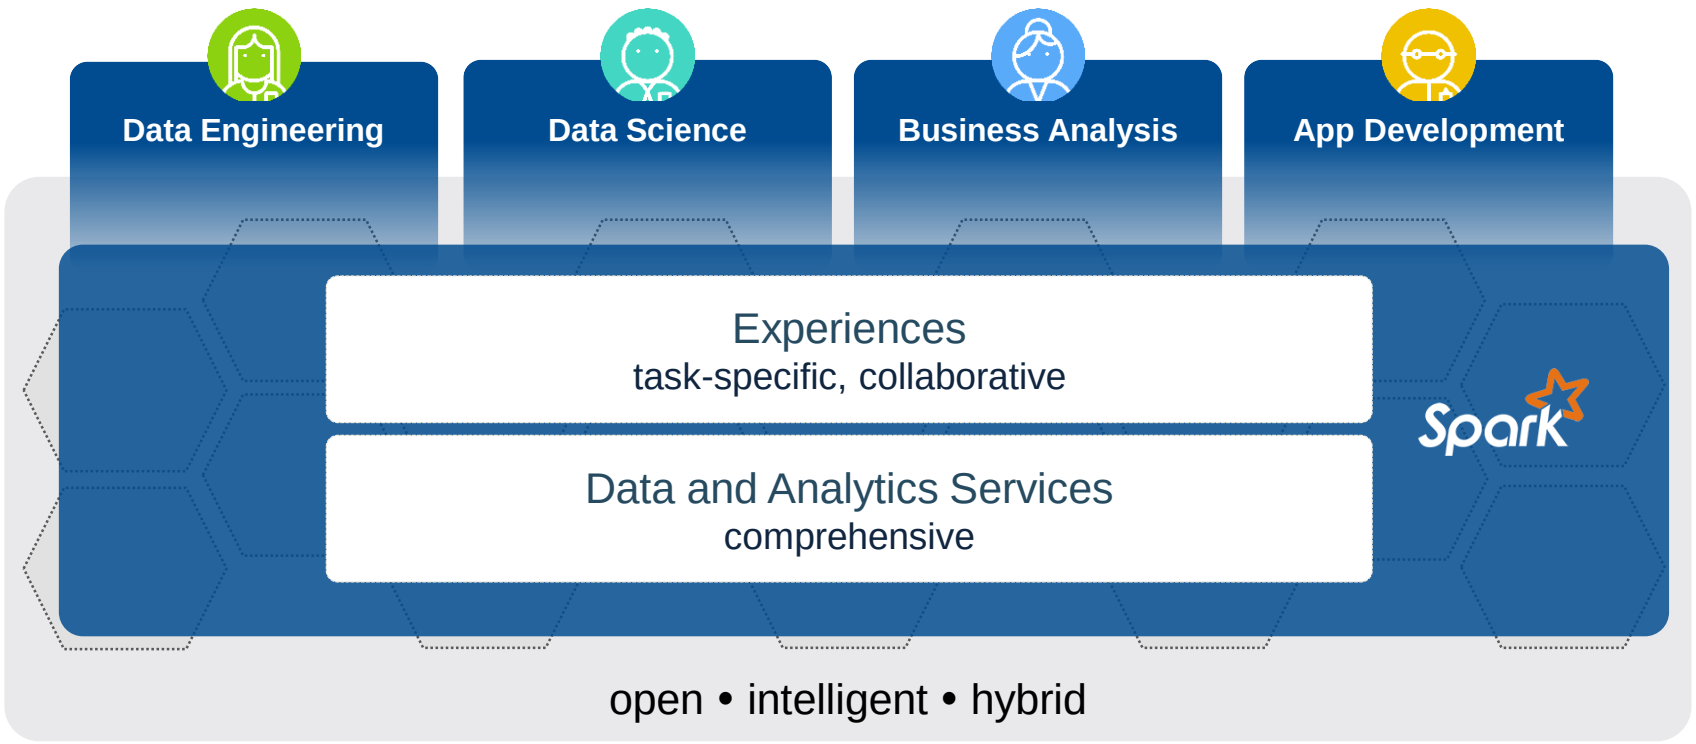

IBM Watson Data Platform

Connects Users to Data and Analytics

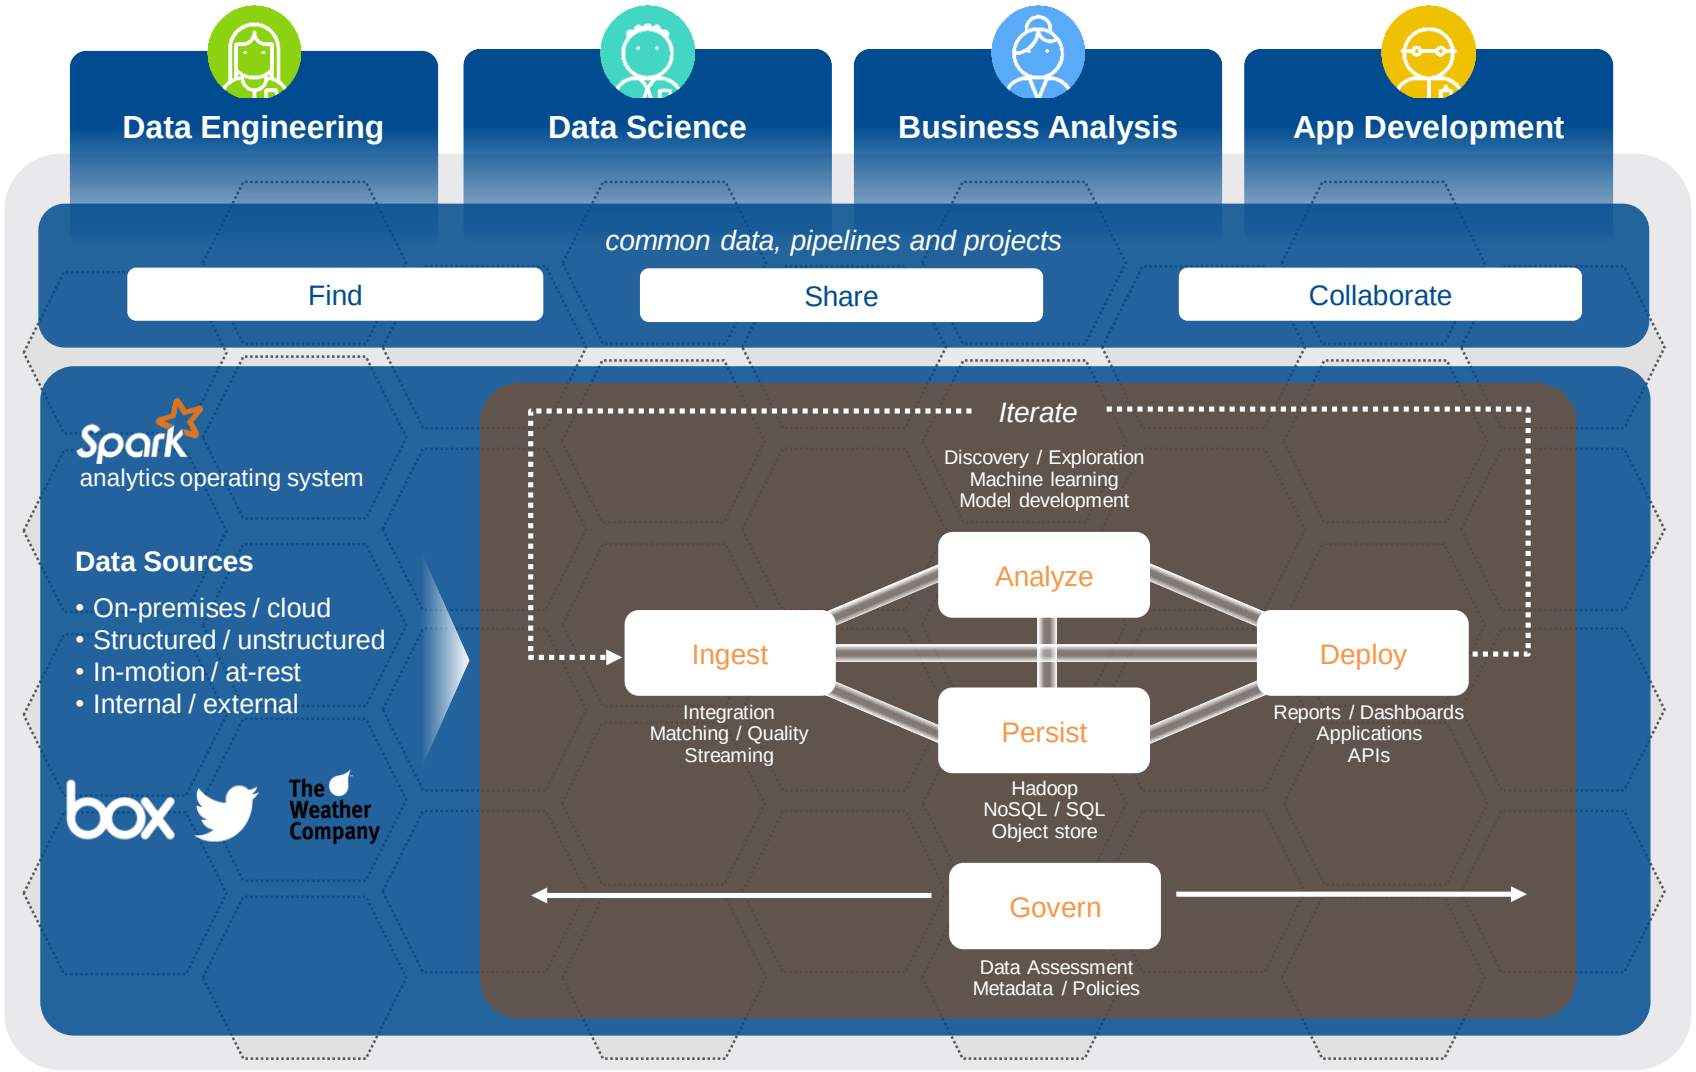

Data Scientist Challenges

- **Rigid toolset**
  - Have to choose one and only one approach
  - Cannot easily connect all of the capabilities needed
  - Difficult to navigate between the various tools used

- **Fragmented and time consuming**
  - Using multiple disjointed environments
  - Separate on-ramp/community for each tool/environment
  - Does not have meta data or data lineage

- **Analytical Silo**
  - Difficult to maintain and version control project assets
  - Limited means of collaborating with team
  - Results are difficult to share

Data Science Experience

Brings together popular Data Science **Open Source tools** with IBM value-add functionalities coupled with **community and social** features

Learn

Built-in learning to get started or go the distance with advanced tutorials

Create

The best of open source and IBM value-add to create state-of-the-art data products

Collaborate

Community and social features that provide meaningful collaboration

External URL: http://datascience.ibm.com

Core Attributes of the Data Science Experience

IBM Data Science Experience

Community

- Find tutorials and datasets
- Read articles and papers
- Connect with Data Scientists
- Share comments
- Copy and share notebooks

Open Source

- Code in Scala/Python/R/SQL
- Jupyter Notebooks
- RStudio IDE and Shiny
- Apache Spark
- Your favorite libraries

IBM Added Value

- IBM Machine Learning*
- SPSS Modeler Canvas*
- Prescriptive Analytics - DOcplexcloud
- Projects and Version Control
- Managed Spark Service

Powered by IBM **Watson Data Platform**

* Closed beta

DSX Architecture

DSX architecture

Last updated: **June 27, 2017**

DSX provides you with the environment and tools to solve your business problems by collaboratively analyzing data. This illustration shows how the architecture of DSX is centered around the project. A project is how you organize your resources for solving a business problem.

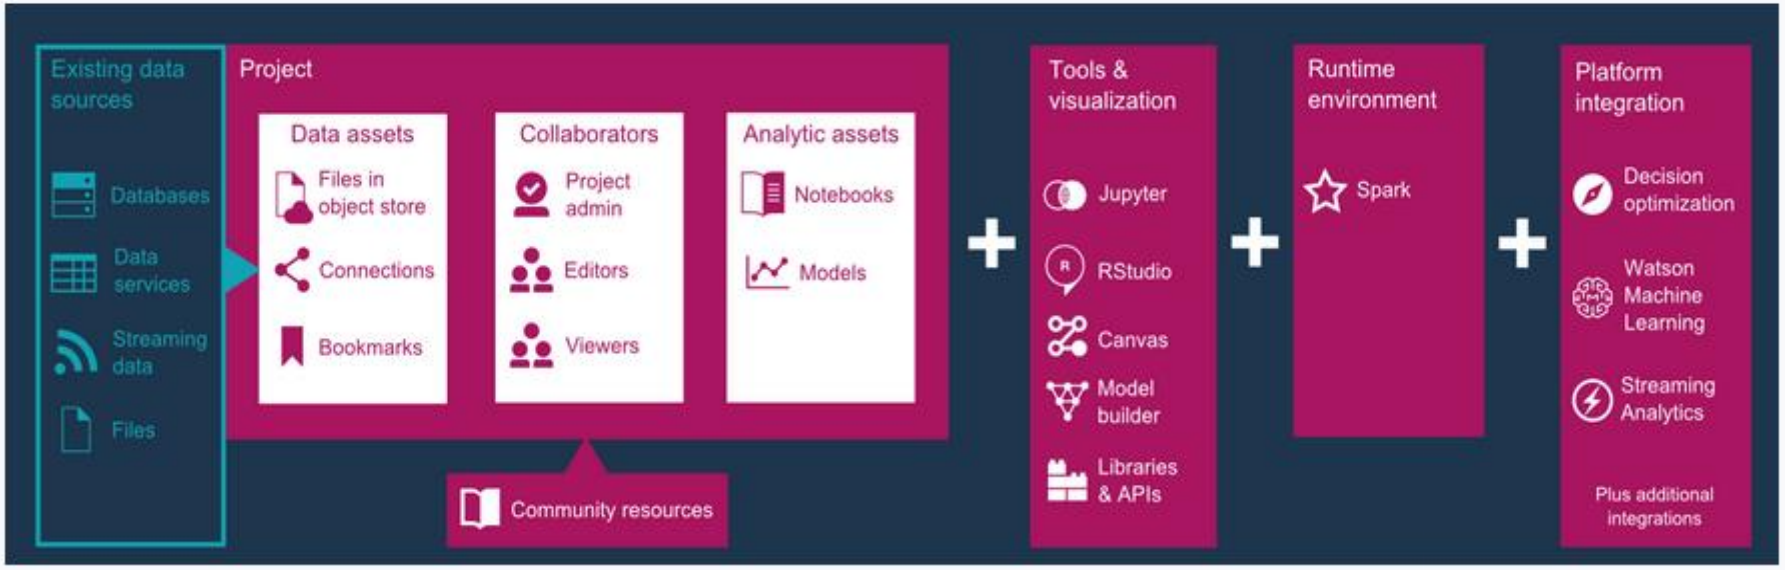

Community Cards provide in-context learning for users

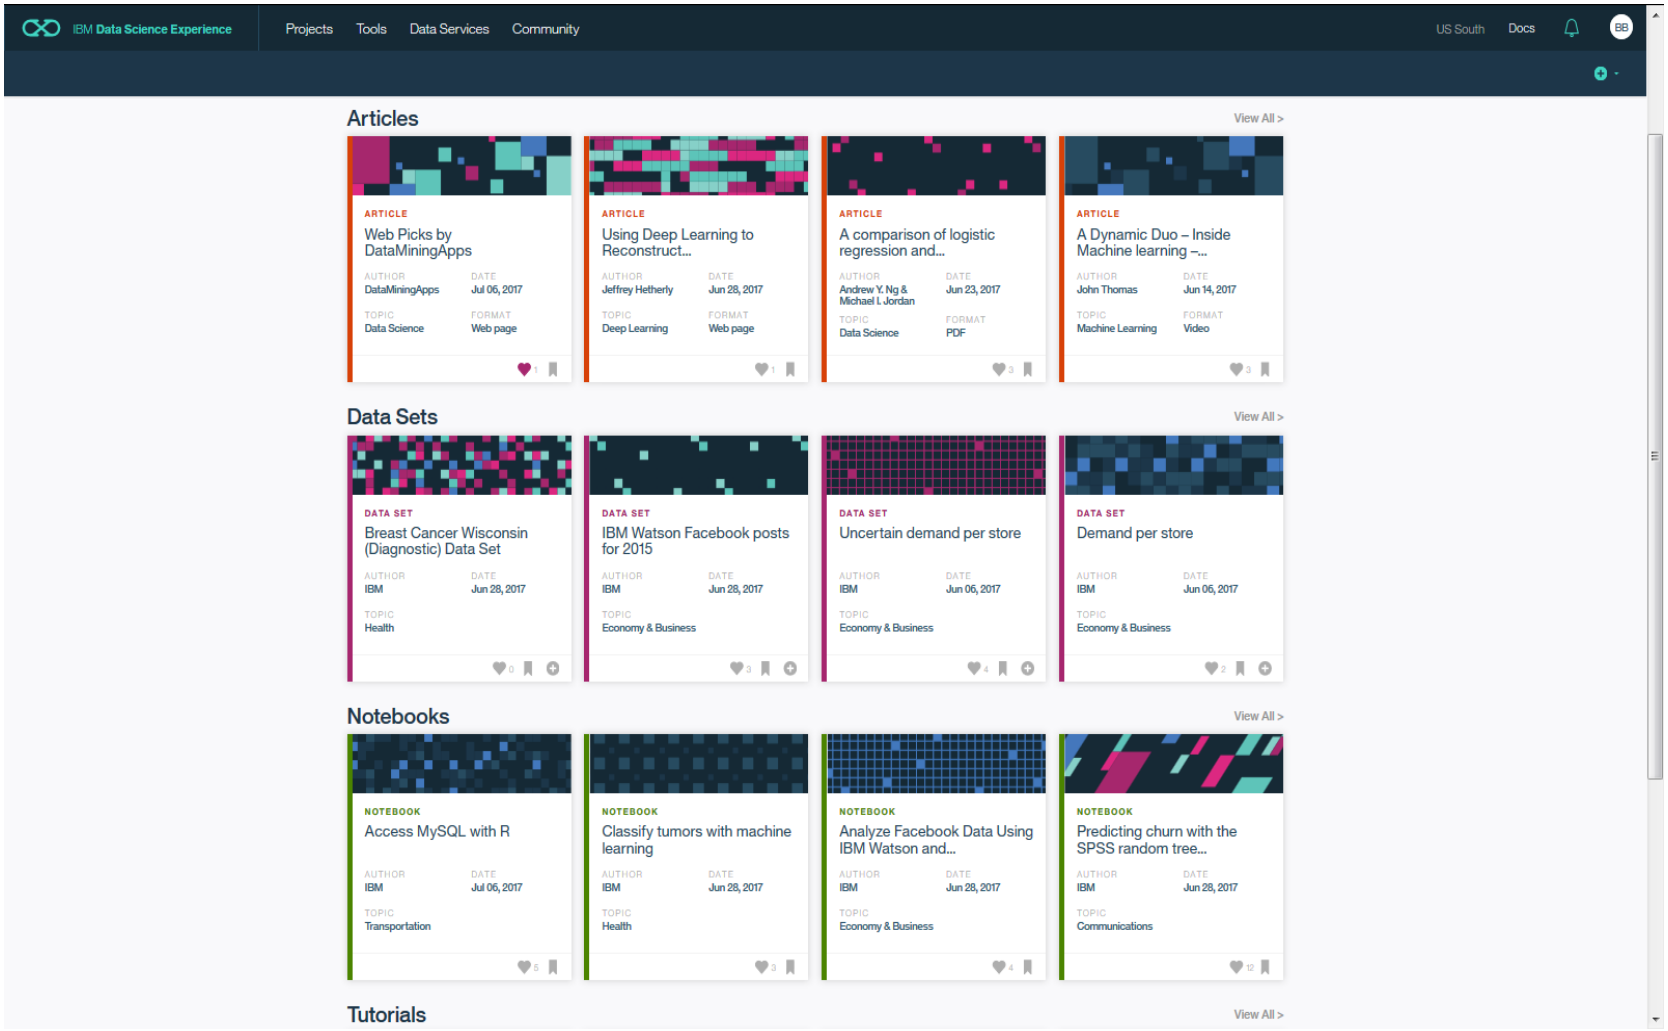

Collaborate Using Projects

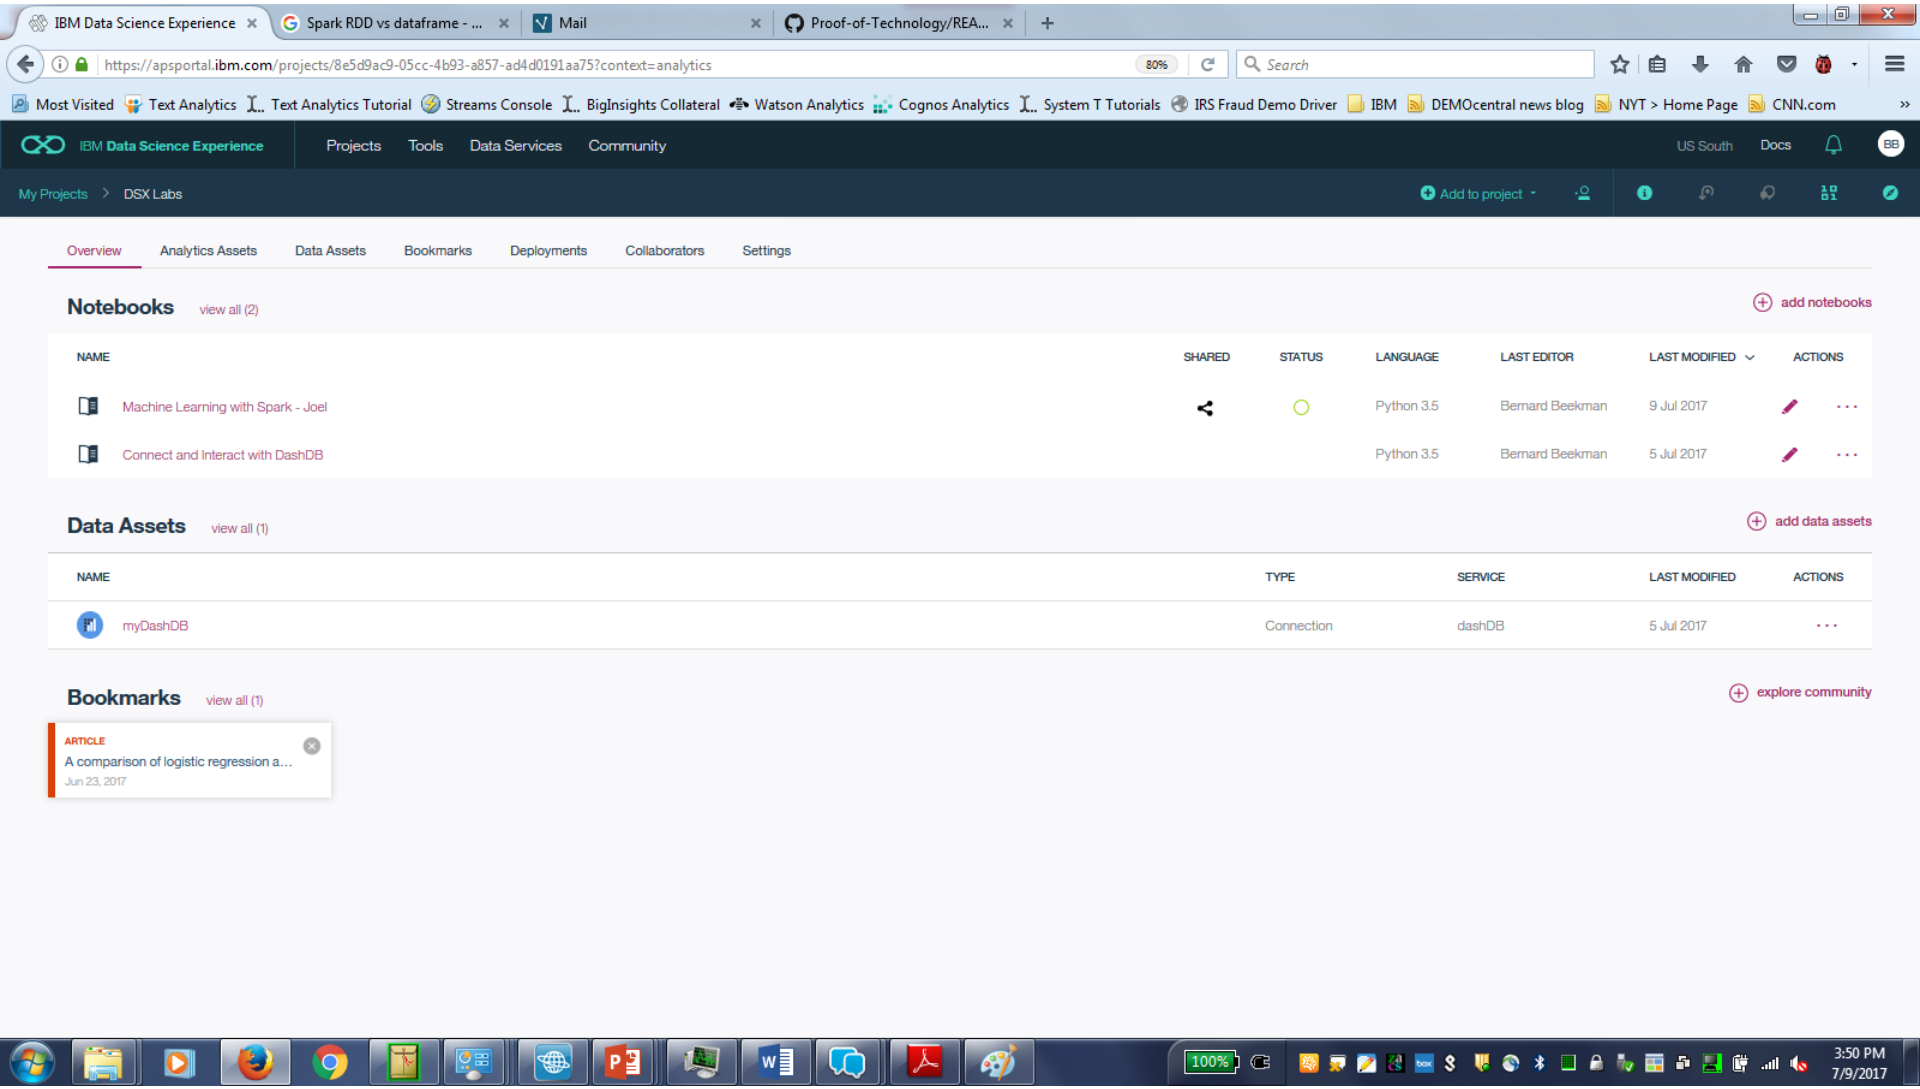

Add Collaborators to a Project

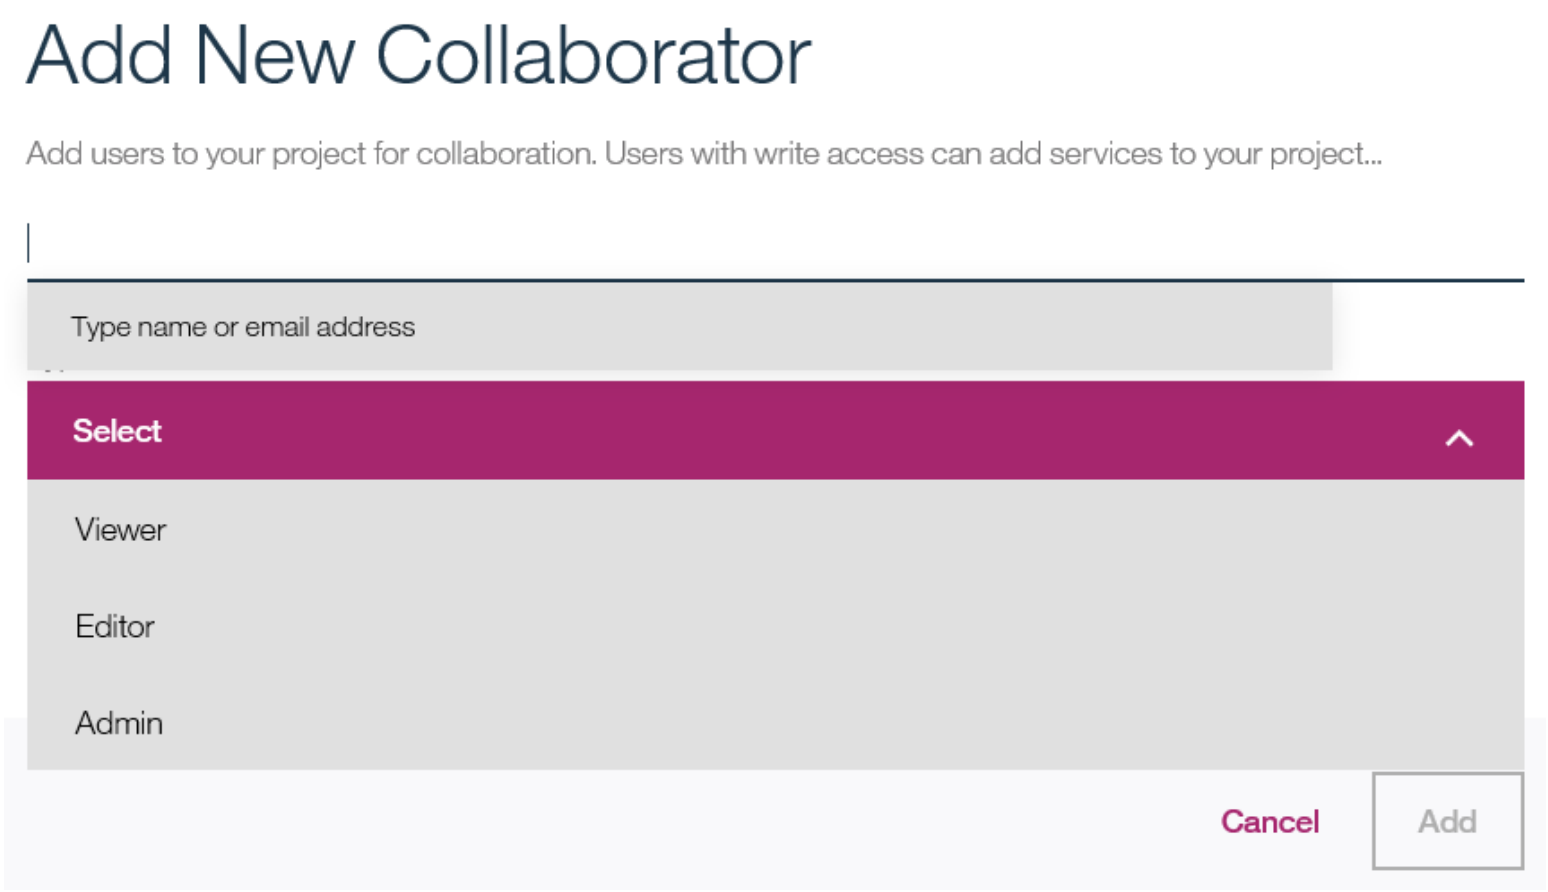

GitHub Integration

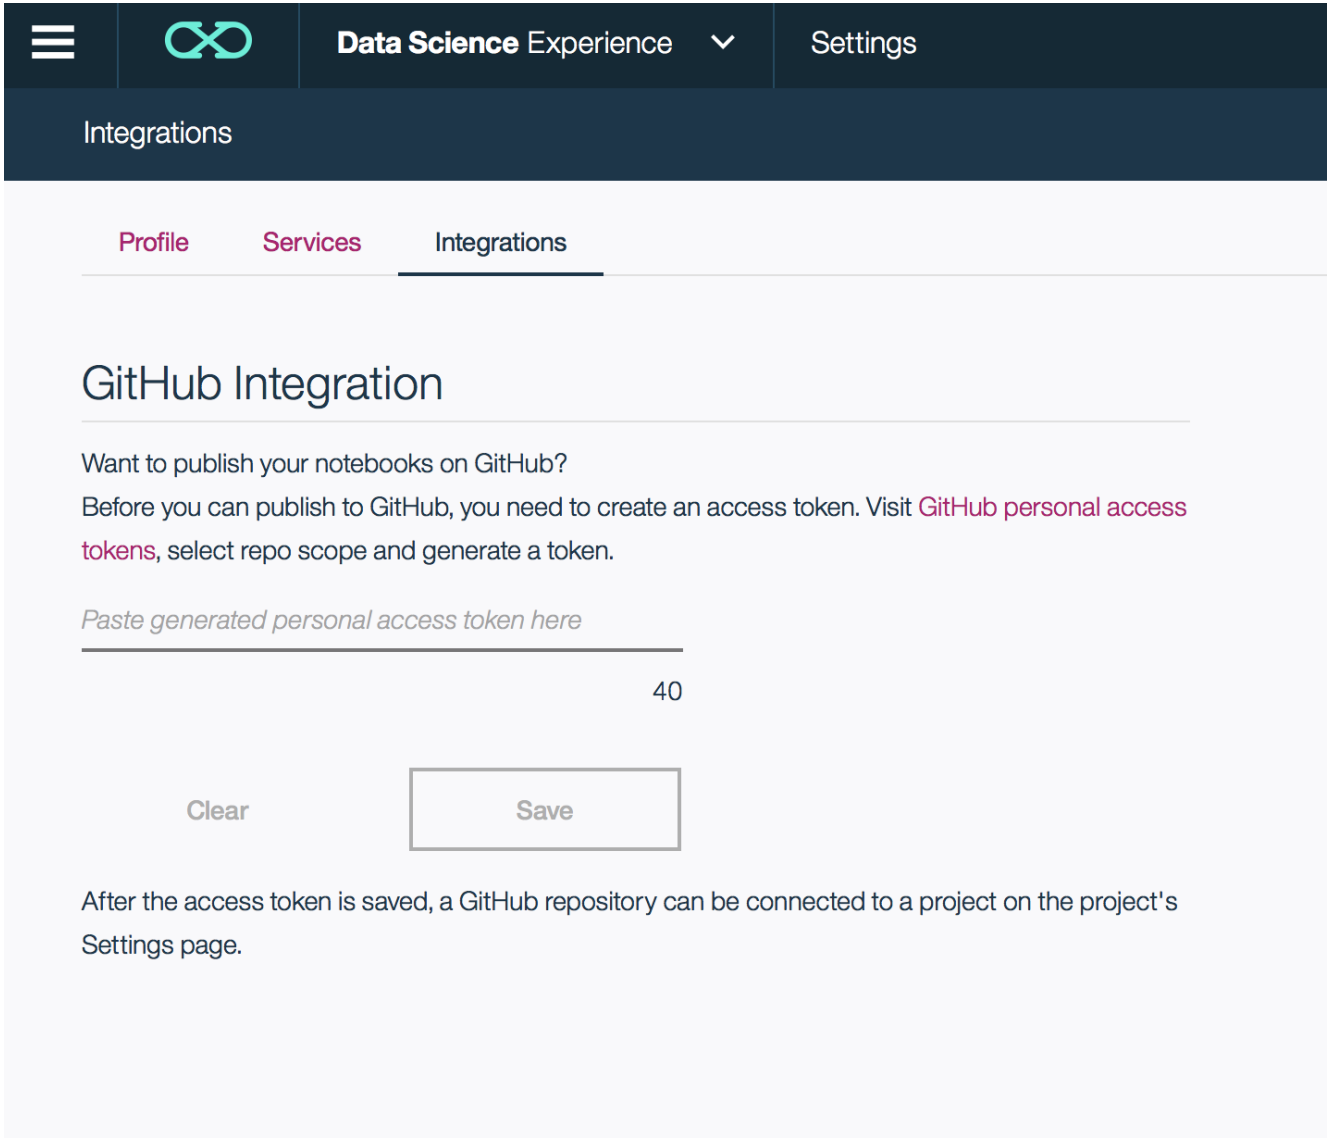

Live chat on Intercom for support from the IBM team and to provide your feedback on how we can improve DSX

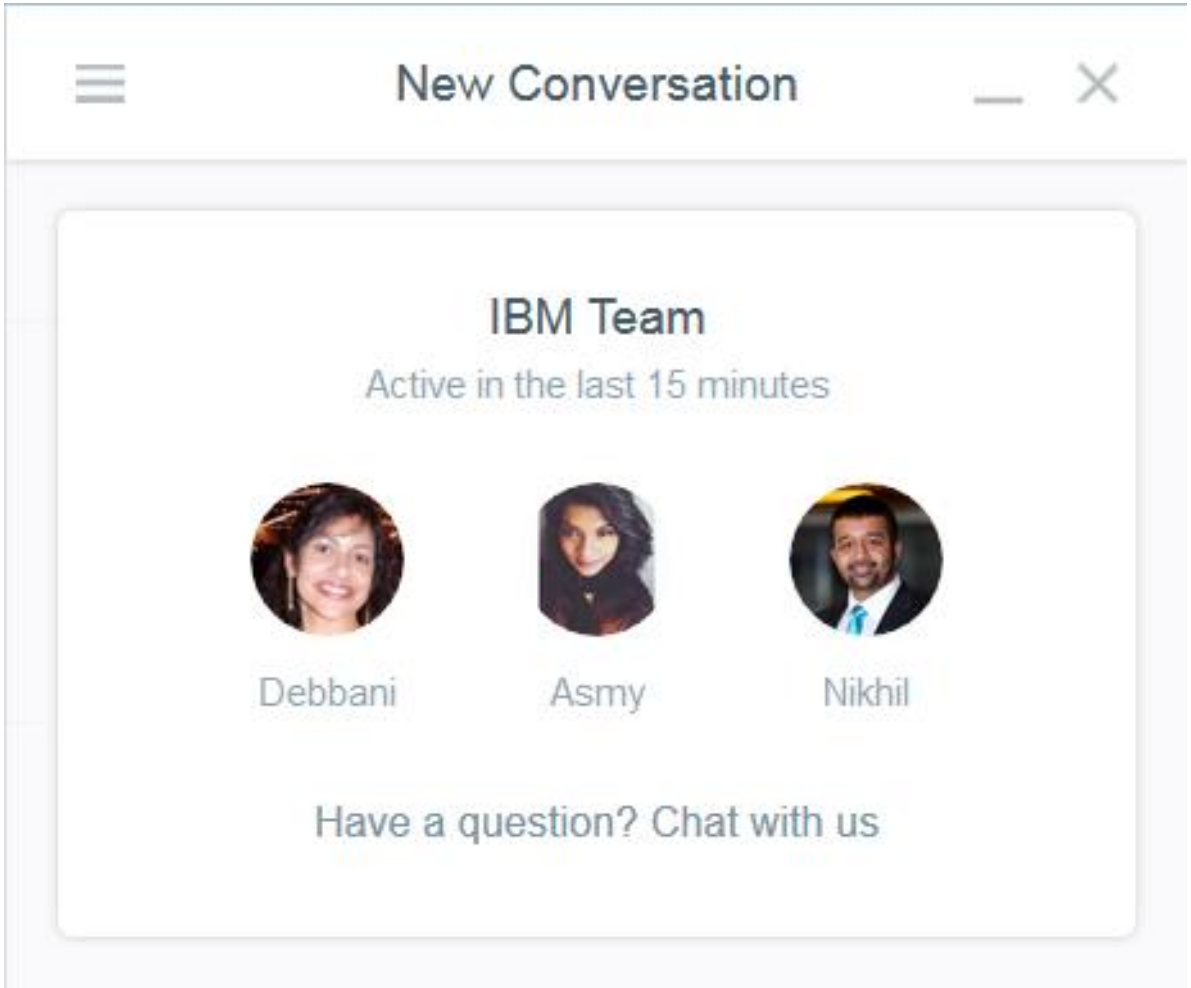

What is a "Notebook"?

Pen and Paper

Pen and paper has long provided the rich experience that scientists need to document progress through notes and drawings:

- Expressive
- Cumulative
- Collaborative

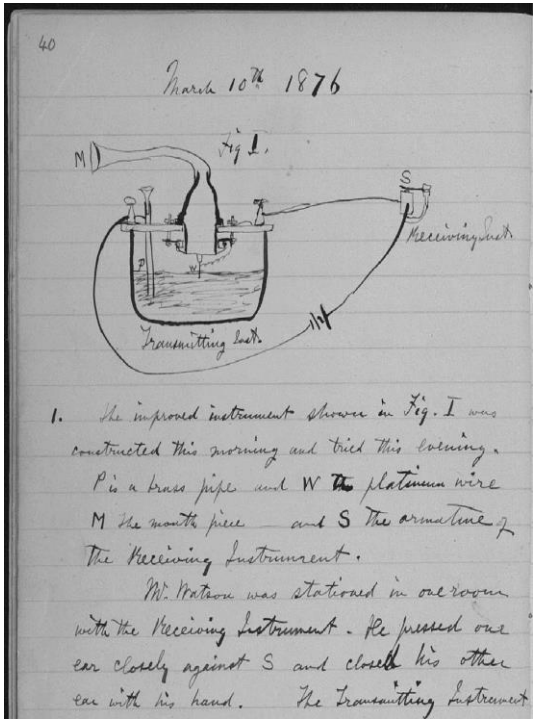

Notebooks

Notebooks are the digital equivalent of the "pen and paper" lab notebook, enabling data scientists to document reproducible analysis:

- Markdown and visualization
- Iterative exploration
- Easy to share

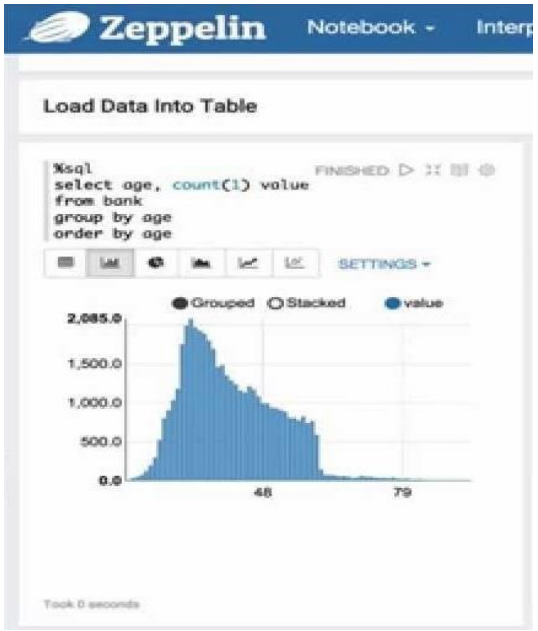

Integrated Jupyter Notebooks for interactive and collaborative development - seamless execution on Spark

Notebooks › Spark Summit West - BlocPower Demo

File Edit View Insert Cell Kernel Help | Python 2

Format Markdown Cell Toolbar None

Explore the correlations

In [9]:
```
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('ggplot')
ff = pd.tools.plotting.scatter_matrix(dfScaled, diagonal='hist',figsize=(10,10))
```

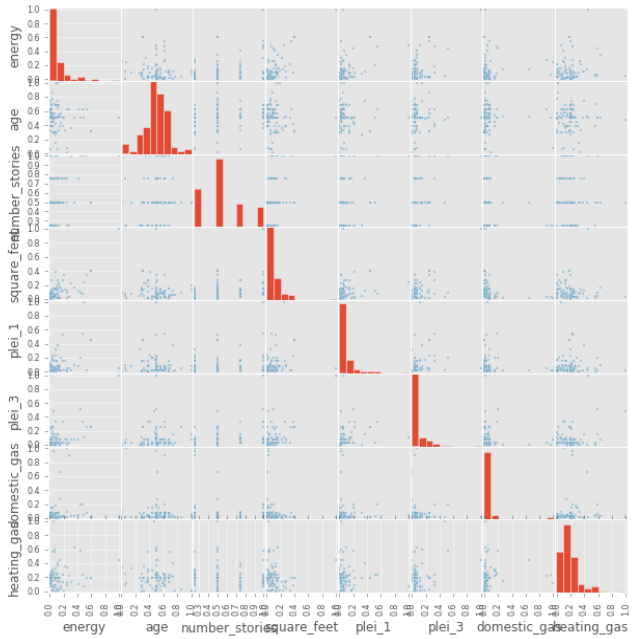

The Spark service uses Bluemix Object Storage as its preferred data store for building performant applications

- Object storage provides inexpensive, scalable and self-healing retention of massive amounts of unstructured data
- Every object exists at the same level in a flat address space
- Bluemix Object Storage has a drag-and-drop upload and Swift API for programmatic access

Object Storage

IBM

Supported Data Sources/Targets for DSX via on- premises and cloud Connectors

Cloud Sources	On-Premises Sources	Cloud Targets	On-Premises Targets
Amazon Redshift	Apache Hive	Amazon S3	IBM DB2® LUW
Amazon S3	Cloudera Impala	Bluemix Object Storage	IBM Pure Data for Analytics®
Apache Hive	IBM DB2® LUW	IBM Cloudant™	Teradata
Bluemix Object Storage	IBM Informix®	IBM dashDB	
IBM BigInsights™ on Cloud *	IBM Pure Data for Analytics®	IBM BigInsights™ on Cloud *	
IBM Cloudant™	Microsoft SQL Server	IBM DB2® on Cloud	
IBM dashDB	MySQL Enterprise Edition	IBM SQL Database	
IBM DB2® on Cloud	Oracle	IBM Watson™ Analytics	
IBM SQL Database	Pivotal Greenplum	PostgreSQL on Compose	
Microsoft Azure	PostgreSQL	SoftLayer Object Storage	
PostgreSQL on Compose	Sybase		
Salesforce	Sybase IQ		
SoftLayer Object Storage	Teradata		

All of the supported targets are compatible with each source

DSX has RStudio built into the experience thanks to our strategic partnership

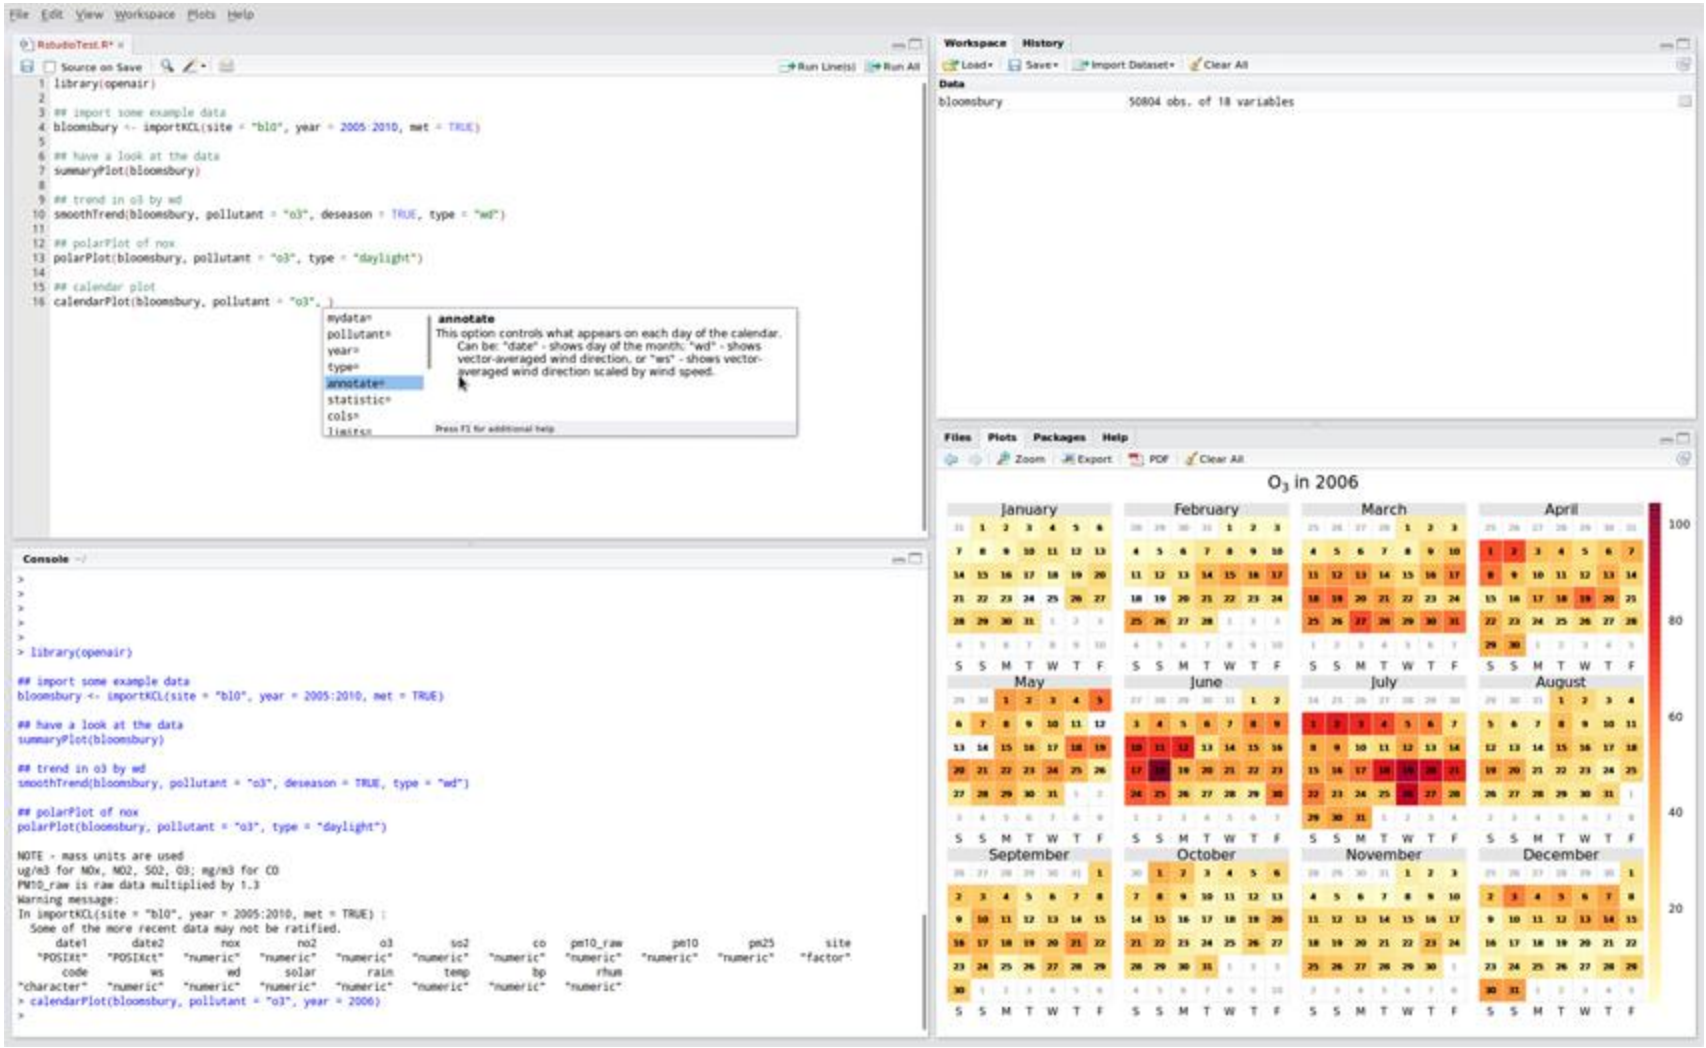

With RStudio you can create Shiny web applications to make your analysis accessible to the business

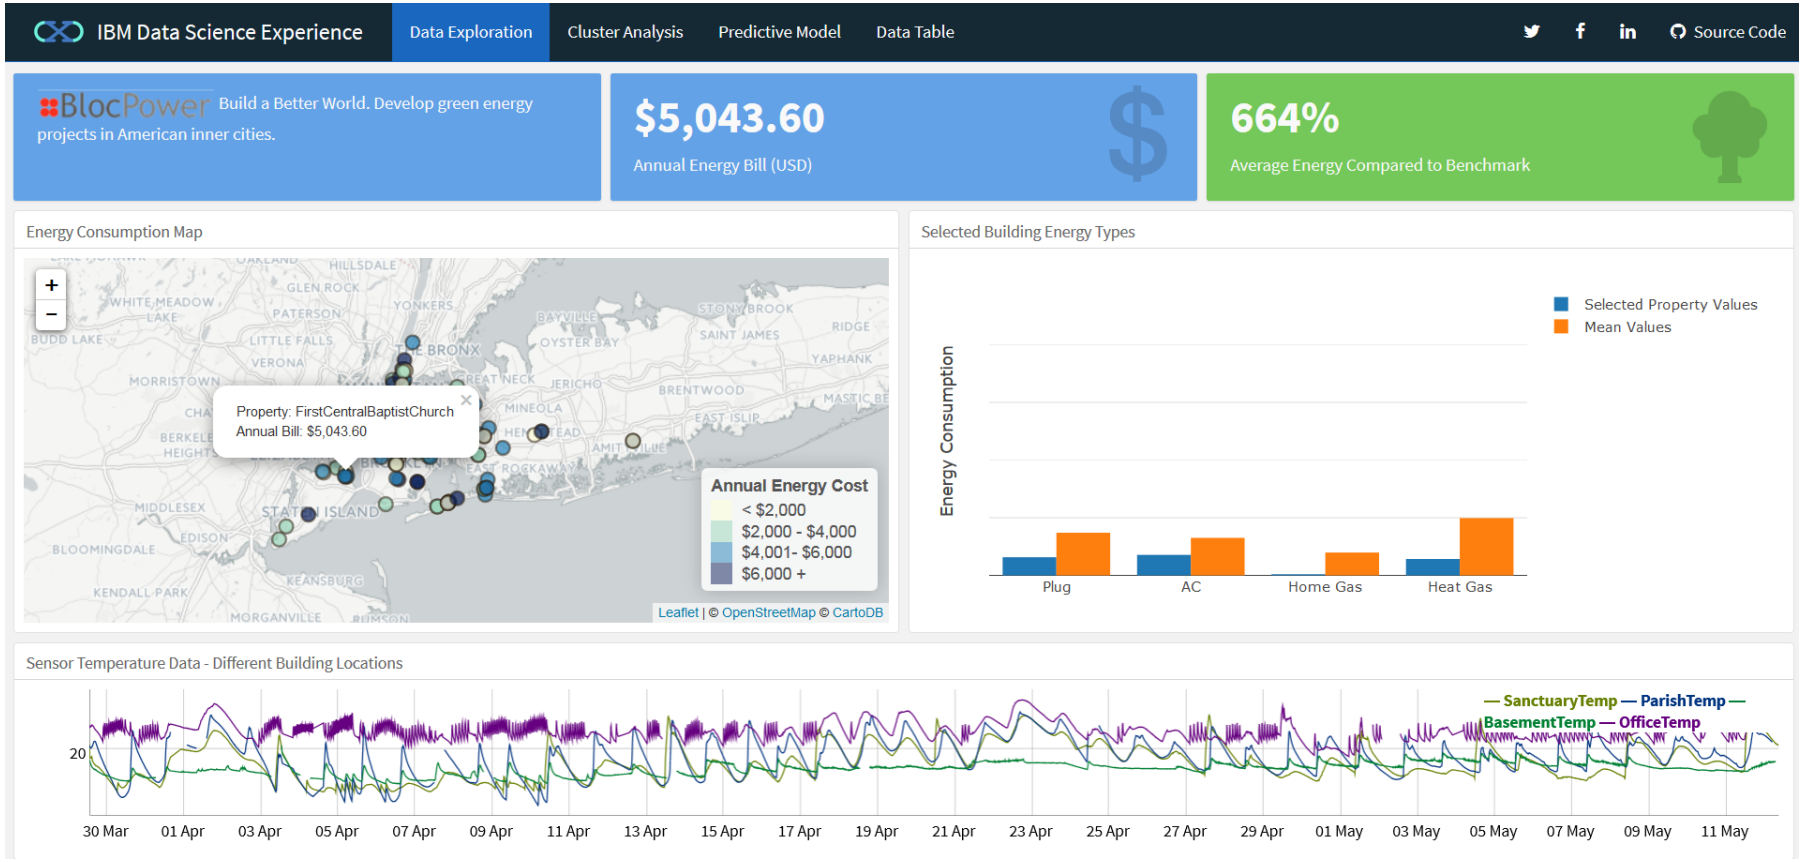

DSX Local

- **Very similar to the public cloud version of DSX**
- **Runs on hardware that is provided by the customer**
  - The DSX Local software and hardware are managed by the customer
- **DSX Local comes with all the software it needs to run, although it can integrate with existing customer systems such as**
  - Databases and HDFS storage
  - LDAP servers for authentication

Get Started with Data Science Experience Today!

Calling all Data Scientists!

- Our mission is to win the **hearts and minds** of Data Scientists
- IBM Data Science Experience is a **freemium model** with value-add features, pricing and up-sell in development
- **Sign up** and encourage your colleagues to do so at **datascience.ibm.com**

IBM Data Science Experience

https://www.youtube.com/watch?v=1HjzkLRdP5k&t=29s